


Two explanations. One separating test.

Measured in DeepMind's Tracr reversal Transformer

Known circuit • frozen rule • 128 fresh token-assignment families



Native first answer: C. The patch copies only the intermediate address (third content position).

Changing donor content separates the explanations; changing recipient content tests the survivor.

128/128 fresh families passed

One-sided 95% lower bound: 97.69% • Frozen adequacy target: 95%

The same classifier also recognized deliberate answer copying in 8/8 development controls.

What makes this more than an output change?

Declared rival predictions; a separating condition chosen from those predictions; audited tensor replacement; paired precision checks; a frozen fresh-family test.

What this establishes

The method reproduces the expected narrowing in an externally built, known circuit. The site was chosen using compiler structure. This validates the implementation and tested candidate comparison; it does not discover an unknown LLM mechanism.